

Contents

Capabilities Mapping 1

Gartner MQ / Critical Capabilities — AI Application Development Platforms 2026 1

1. The rules that shape everything 1

1.1 Cutoff — preview features are not evaluated at all 1

1.2 Pro-code is the posture 2

1.3 Mandatory features — all seven must be visible on camera 2

1.4 Optional features — claim only what is real 2

2. Use cases 3

3. Critical Capabilities — the scoring spine 3

4. Requirement coverage 3

4.1 Section 1 — deep research agent (reqs 1.1–1.31) 3

4.2 Section 2 — external customer assistant (reqs 2.1–2.11) 4

4.3 Section 3 — customer agent, grounding, guardrails, DLP, reports (reqs 3.1–3.15) 4

4.4 Section 4 — avatar / digital human (reqs 4.1–4.8) 4

5. Regional and regulatory positioning 4

6. Verification before recording 5

6.1 Probe results — account 157000088284, us-east-1, 3 Aug 2026 5

7. Known gaps and how to handle them 6

Capabilities Mapping

Gartner MQ / Critical Capabilities — AI Application Development Platforms 2026

Decision record for what we claim, where we prove it, and what must be verified first. Timecode index lives in `table-of-contents.md`; talk tracks in `demo-script.md`.

Sources (in docs/, git-ignored — Gartner-RESTRICTED, never commit): `Welcome Packet for Magic Quadrant.pdf` · `Magic-Quadrant-and-Critical-Capabilities.xlsx` (the RFI) · `Vendor Demonstration Guidelines ... 2026.pdf` — **not yet in this repo; add before using §4 as a checklist**.

Status legend — GA confirmed generally available · GA-VERIFY believed GA, confirm against the 1 Aug 2026 service state (§6) · PREVIEW must not be presented as GA · 3P predominantly third-party, disclosed.

Two taxonomies — do not conflate. The Welcome Packet defines twelve Critical Capabilities and the published report scores against those. The RFI's Subsection Name column tags only ten: Agent State Management is filed under *Agentic Framework* (Q35), Deployment is split across *Model Catalog* (Q28) and *Orchestration* (Q43–47). §3 is keyed to the Welcome Packet's twelve. The RFI column tells you where to type the answer.

1. The rules that shape everything

1.1 Cutoff — preview features are not evaluated at all

"Capabilities must be in production (or generally available) as of August 1, 2026... We will not evaluate products or capabilities that are available as a limited release or beta version."

A preview feature shown as GA is not discounted, it is **discarded** — and it casts doubt on every other claim. Excluded from the video, declared as beta in the RFI (the dropdown has a distinct "Yes, in beta today" option):

Capability	Status	Handling
AgentCore Registry	Preview (Apr 2026)	Cut from the video in the retime to 52:30 — a preview capability is not evaluated, so the seconds went elsewhere
AgentCore Payments	Preview (May 2026)	Not in the video
Optimization → Failure Insights	Preview (Jun 2026)	Omit. Recommendations, Config Bundles, A/B are GA
Harness → Claude Agent SDK export	Coming soon	Only Strands export exists
AgentCore CLI → Terraform	Coming soon	Claim CDK only

1.2 Pro-code is the posture

“We exclude vendors who emphasize no-code... or have a go-to-market strategy focused on low-code application platforms or primarily serving business users.”

Inclusion additionally requires pro-code capabilities, **software engineers as a core user persona**, and building without mandated professional services. RFI Q5 is multi-select across Low/No/Pro-Code — answer all three there. But the **video's centre of gravity is the software engineer**: Harness framed as configuration-as-code in git, not a canvas; Marcus owns Blocks 1A–1G; Lena under three minutes as a reviewer inside an engineer-built interface. Enforced by V15.

1.3 Mandatory features — all seven must be visible on camera

Miss one and the market definition itself is in question.

Mandatory feature	Where
Pro-code framework for assistants, agents, multimodal apps	Kiro specs; Harness as config-as-code; Strands export; any framework on Runtime
Grounding on organizational knowledge	Bedrock Managed KB (S3/SharePoint/Confluence/Drive/OneDrive/Web Crawler); context engineering
Guardrails in and out	Bedrock Guardrails at model, at gateway (Guardrails in Policy), and on voice
Model catalog, commercial and open-source	Bedrock catalog, model cards, AI Service Cards; Mantle; LiteLLM
Deployment — cloud and hybrid	Runtime; plus EKS and on-prem Helm over PrivateLink ; Evaluations/Optimization govern agents on Lambda, EKS
Evaluations, prebuilt or custom	AgentCore Evaluations — built-in, custom LLM, custom code (Lambda)
Orchestration for agent-to-agent collaboration	Supervisor + parallel section agents; A2A with agent cards; Gateway HTTP targets; Step Functions

Hybrid was the easiest of these to lose. It is mandatory, it is a live beat at 25:30, and V16 guards it.

1.4 Optional features — claim only what is real

Optional feature	Answer
Security & risk (DLP, sandbox, IAM)	Guardrails sensitive-information filters; Code Interpreter Sandbox network mode; Identity; Policy; PrivateLink
AI router — route by use case, performance, cost	Gateway inference targets — inference connector/provider targets, model discovery, ID translation, per-model token lin
Observability across the stack	AgentCore Observability / CloudWatch GenAI Observability, OTEL-compatible
Advanced grounding — knowledge graph, chunking, rerank	Managed KB hybrid search, chunking, ranking/reranking. Graph search unresolved — V6.
Composability	Any framework; any model; OTEL to third-party; MCP and A2A; Cedar
Catalogs & marketplaces	Gateway connectors and integration templates; AWS-curated skills catalog; AWS Marketplace AI Agents and Tools; Bedro

2. Use cases

Three application types, all three required by the inclusion criteria. Timecodes and spans in `table-of-contents.md` §1.

Use case	Carried by
AI assistants	External customer assistant grounded on the step-1 report — text, images, voice; automatic quality control; verified page-level grounding and correct refusals
AI agents	Deep research agent (plan → parallel multi-agent → synthesis) and the account-opening agent (KYC, eligibility, fraud, auditable decline)
Multimodal applications	Multimodal input (regulatory PDF, chart image, agent reasoning over its own generated chart) and output (logo, storefront, service tiles, speech); one chat

3. Critical Capabilities — the scoring spine

All twelve from the Welcome Packet. Timecodes in `table-of-contents.md` §3.

#	Capability	RFI Qs	Primary AWS answer
1	Grounding	4, 6, 7	Bedrock Managed KB — hybrid search, chunking, encoding, ranking, reranking, six native connectors
2	Guardrails	8–13	Bedrock Guardrails — content filters, denied topics, word filters, prompt-attack detection, sensitive topics
3	Evaluations	14–20	AgentCore Evaluations — built-in evaluators at session/trace/tool level, custom LLM and code-based
4	Observability	21–24	AgentCore Observability / CloudWatch GenAI Observability — span trees, per-span latency/token
5	Model Catalog	25–31	Bedrock catalog, model cards, AI Service Cards; Amazon Nova family; customization, fine-tuning, c
6	Agentic Framework	32–36	AgentCore Harness, Runtime, Memory; Strands Agents (first-party open source); prebuilt tool integ
7	Multimodal Framework	37–39	Amazon Nova multimodal models; image generation; speech-to-speech; Bedrock Data Automation;
8	Cost Management	40–42	Per-session cost by model and agent; allocation tags; budget alarms; prompt caching; batch inferenc
9	Agent State Management	35 (→ <i>Agentic Framework</i>)	Continuity: Memory short/long-term strategies, cross-session context, session filesystem surviving
10	Deployment	28, 43–47 (→ <i>Model Catalog / Orchestration</i>)	Runtime immutable versions and named endpoints with instant rollback; EKS and on-prem Helm
11	Orchestration	43–47	Gateway as AI gateway — token tracking/reporting, rate limiting, user limits, response caching
12	Security	48–52	AgentCore Identity (3LO, OBO, Private Key JWT, token vault with CMK); Policy (Cedar + automa

Rows 9 and 10 are the two most likely to be under-served, because the RFI hides them under other tags and both are easy to assert rather than show:

- **Agent State Management** — the memory half demos itself; the operational half does not. Beat 17:00 must actually fail a tool mid-pipeline and show detection → retry → resume with state intact. Claiming stuck detection while nothing gets stuck is the weakest form of coverage.
- **Deployment** — hybrid is a mandatory feature (§1.3) and Gartner separately asked for deployable containers or Helm. A slide does not satisfy it. See V16.

Combinations. Gartner asks specifically to “showcase how you combine various capabilities.” Six are staged deliberately — governed research · auditable multi-agent · layered defence · quality as a release gate · closed improvement loop · one corpus four channels. Definitions and timecodes in `table-of-contents.md` §3.1. Name each one on camera and in the RFI.

4. Requirement coverage

Every requirement in the demo brief, grouped. **All GA unless flagged.**

4.1 Section 1 — deep research agent (reqs 1.1–1.31)

Reqs	Requirement group
1.1–1.6	Full lifecycle: build with an agentic framework, design, implement, test, deploy, operationalize
1.7–1.11, 1.13–1.14	Refine the query, decompose into sub-questions, produce a research plan naming objectives / methods / evaluation criteria / expected outcomes, generate search
1.12	Tool calling / MCP such as web search APIs
1.15–1.16, 1.18	Multimodal input incl. data-analysis charts; multimodal output incl. graphics/logos/storefronts; one chapter as text + images + speech
1.17, 1.19–1.20	Combine sections with introduction and conclusion; citations for all sources; export as PDF
1.21–1.31	IDE / low-code / high-code / graphical development · model cataloguing · data preparation · tracing · evaluation · debugging · agent operations · observability

4.2 Section 2 — external customer assistant (reqs 2.1–2.11)

Reqs	Requirement group	AWS answer
2.1–2.2	Assistant for external customers grounded on step 1; detailed service info as text and images	Managed KB over the step-1 PDF and imagery; BDA parses imagery so retrieval
2.3–2.4	Generate service images; speak account and investment progress	Nova image generation on demand within brand and content guardrails; Nova
2.5–2.8	Automatic quality control — format, length, filter, irrelevant information	Response schema validator (mandatory fee disclosure); channel length policy
2.9	Human review to update service descriptions	Gateway elicitation Form mode; versioned content with attribution and diff
2.10–2.11	A/B test output across various models; continuous feedback loop	Optimization A/B — config-bundle variants, gateway traffic split, statistical

4.3 Section 3 — customer agent, grounding, guardrails, DLP, reports (reqs 3.1–3.15)

Reqs	Requirement group	AWS answer
3.1–3.4	Open an account; validate you can do business with them; take advantage of services; detect fraudulent accounts and transactions	Multi-step agent with visible tool calls to
3.5–3.7	Grounding via context/prompt engineering and RAG; quickly verify answers are grounded in the step-1 PDF	Inspectable structured context object with
3.8–3.10	Guardrails block non-bank requests; employee-salary questions trigger guardrails; DLP measures fire	Denied topics and content filters; entitlement
3.11–3.15	Out-of-the-box reports: comprehensiveness, accuracy, response time, costs, and metrics for marketing	Completeness by intent (trended); Group

4.4 Section 4 — avatar / digital human (reqs 4.1–4.8)

Reqs	Requirement group	AWS answer
4.1	Photorealistic avatar	Partner rendering layer on the AWS voice stream — 30-second disclosure, not demon
4.2, 4.5	At least two languages; guidance in multiple languages and voice tones	Nova speech-to-speech with native per-language grounding; prosody and tone control
4.3–4.4	Ground the avatar to step-2 text and images; spoken Q&A on why to choose the bank	Same Managed KB and content objects as the Block 2 assistant; live grounded voice
4.6–4.7	Show / indicate the item being discussed; change item on customer question	Agent-emitted UI tool call highlighting the tile, visible in the trace; mid-answer topic
4.8	Test guardrails with out-of-scope questions	Three voice probes blocked by the identical rule set; unified audit log

5. Regional and regulatory positioning

The scenario names US, Europe and China. The Welcome Packet also lists China among the geographies for the 10+ enterprise customer test — so China needs a precise answer, not a smooth one.

Requirement	Position
EU data residency	Cross-Region inference profile scoped to EU Regions; KB and Runtime in-Region; EU-resident reviewer operating in the EU plane
EU suitability / MiFID II	Spec requirement 4 traces to the MiFID II extract dropped in at 01:45; denied topics block unlicensed advice
US regulatory	KYC and sanctions screening, entitlement-aware retrieval, auditable decline reasons, retained decision records
Network isolation	Every AgentCore capability supports VPC and PrivateLink; agent traffic need not traverse the public internet
Attestation	SOC 1 / 2 / 3 (Jun 2026); GovCloud (US-West)
China	State plainly: AWS China Regions are a separate partition under a local operating entity . Trinity's China business runs as a separate deployment

Do **not** imply a single global control plane spanning the China partition. Framing it as a repeatable separate deployment is both accurate and a strength — the packaging deliverable is what makes it repeatable.

6. Verification before recording

Every row is a claim that appears on camera. Anything unconfirmed is cut from the video and declared in the RFI instead. **V-numbers are referenced by demo-script.md and table-of-contents.md — do not renumber.**

ID	Verify
V1	Anything shipped in August 2026 — release notes reviewed for this script end at July 2026
V2	Supported language list for Nova speech-to-speech; pick the second demo language from it
V3	Whether a first-party photorealistic avatar is GA on 1 Aug 2026
V4	GA status of every AgentCore capability named on camera — especially Gateway elicitation, inference targets, HTTP passthrough targets, interactive shells, unified spa
V5	Exact model IDs and Regions , each GA in the demo Region
V6	Graph search / knowledge graph grounding availability (RFI Q4)
V7	Evaluator names on screen match the service. Resolved against the account: 16 built-in evaluators — see §6.1 for the exact list. Name specific evaluators; if you st
V8	Per-unit pricing before quoting any figure. Verified only: Web Search \$7.00/1,000 queries; Identity \$0.010/1,000 requests (free via Runtime/Gateway); Observability = CloudW
V9	Guardrails feature names : content filters, denied topics, word filters, sensitive-information filters, contextual grounding, Automated Reasoning checks, prompt-attack detection
V10	AgentCore Region list before quoting a count
V11	Kiro property-based testing, hooks, Supervised mode and the spec workflow are GA in the build recorded
V12	The China statement in §5, with legal and compliance
V13	How many AWS products Gartner has under evaluation — the guidelines require one demo per product
V14	Source package contains no internal-only references, package names or wiki links
V15	Watch the assembled video against §1.2 — does it read as pro-code for software engineers? Recut Lena if hesitant
V16	The 25:30 hybrid beat shows on-prem/EKS actually working , not a slide
V17	Cross-check video claims against RFI answers before both go in
V18	Web Search Tool Region availability — documented as us-east-1 only , which is the demo Region, so there is no conflict. Reconcile with the EU residency claim at 06:30. S
V19	Web Search acceptable use — citations and links must be retained and displayed in end-user output; no bulk extraction or competing index
V20	Policy in AgentCore GA status and Regions , specifically
V21	Produce and re-verify table-of-contents.md timecodes against the final edit

6.1 Probe results — account 157000088284, us-east-1, 3 Aug 2026

Read-only **List*** calls against the demo account, via boto3 **1.43.63** (the version installed in the project venv). Note the AWS CLI shipped on the build machine — 2.27.35 — does not recognise **bedrock-agentcore-control** at all, so probe through the venv, not the CLI.

Every capability the script depends on is reachable. All 14 probes returned success:

Capability

Harness · Policy engine · Evaluations (evaluators, online config, batch) · Registry · Gateway · Runtime · Memory · Browser · Code Interpreter · Workload Identity · A/B Test

V7 — 16 built-in evaluators, exactly as returned by `ListEvaluators`. The Mar 2026 release note's "13" is stale; use this list:

<code>Builtin.Coherence</code>	<code>Builtin.Conciseness</code>	<code>Builtin.Correctness</code>
<code>Builtin.Faithfulness</code>	<code>Builtin.GoalSuccessRate</code>	<code>Builtin.Harmfulness</code>
<code>Builtin.Helpfulness</code>	<code>Builtin.InstructionFollowing</code>	<code>Builtin.Refusal</code>
<code>Builtin.ResponseRelevance</code>	<code>Builtin.Stereotyping</code>	<code>Builtin.ToolParameterAccuracy</code>
<code>Builtin.ToolSelectionAccuracy</code>	<code>Builtin.TrajectoryAnyOrderMatch</code>	
<code>Builtin.TrajectoryExactOrderMatch</code>	<code>Builtin.TrajectoryInOrderMatch</code>	

Two groups are worth naming explicitly on camera because they answer requirements directly rather than generically:

- **GoalSuccessRate** is the end-to-end task-completion metric the Evaluations capability definition calls out.
- **ToolSelectionAccuracy**, **ToolParameterAccuracy** and the three **Trajectory*Match** evaluators are the *expected tool execution sequence* ground-truth ask. Naming these beats saying "we support tool evaluation."

V5 — all four configured models are accessible. Note the difference in how they resolve, because the residency beat depends on it:

Model	Resolves as
<code>us.anthropic.claude-sonnet-4-6</code>	inference profile (cross-Region)
<code>amazon.nova-2-sonic-v1:0</code>	foundation model
<code>amazon.nova-2-lite-v1:0</code>	foundation model
<code>amazon.nova-2-multimodal-embeddings-v1:0</code>	foundation model

119 foundation models and 63 inference profiles are visible in the account — useful concrete numbers for the model catalog beat at 06:30, though re-check them on the recording day.

V18 — not resolvable by probe. The only connector operations on the control plane are ***PaymentConnector** (the preview Payments feature). Gateway *target* connectors like `web-search` are supplied as `connectorId` at `CreateGatewayTarget` time and there is no API that lists offered connectors. Confirming Web Search therefore requires an actual target creation — do it as a small isolated write, not buried inside a full stack deploy, so a failure is unambiguous.

V20 — partially resolved. `ListPolicyEngines` answers in us-east-1, so the API is live and callable. That is necessary but not sufficient: a callable API can still be public preview. Confirm GA status with the service team before Policy appears on camera, because beats 13:25 and 41:25 and differentiator D2 all rest on it.

Still open and not addressed by this probe: **V6** (graph search), **V10** (Region list), **V8** (pricing).

7. Known gaps and how to handle them

Honesty scores better than coverage. The RFI has explicit "Yes, in beta today" and "Not supported" options, and the source-code deliverable makes overclaims discoverable.

Gap

Photorealistic avatar rendering is partner technology
Registry, Payments, Failure Insights are preview
Knowledge graph / graph search

Gap

AgentCore Web Search is claimed but not wired — `gateway/tools/web_search/handler.py` is a Lambda target calling Nova Web Grounding (`systemTool: nova_grounding`), not Nova Web Search

Terraform for AgentCore CLI; Claude Agent SDK harness export

China partition service and model availability differs

Red team / purple team testing services (RFI Q19)

Sound-effect and music generation models (RFI Q30)

Anything failing §6
